

EDA & Working with Text

(Reading: 5.3-5.7, 8)

Learning goals:

- Understand the goals of EDA: Data Types, Granularity, Scope, Temporality, and Faithfulness
- Learn and practice using regular expressions

UC Berkeley Data 100 Summer 2019
Sam Lau

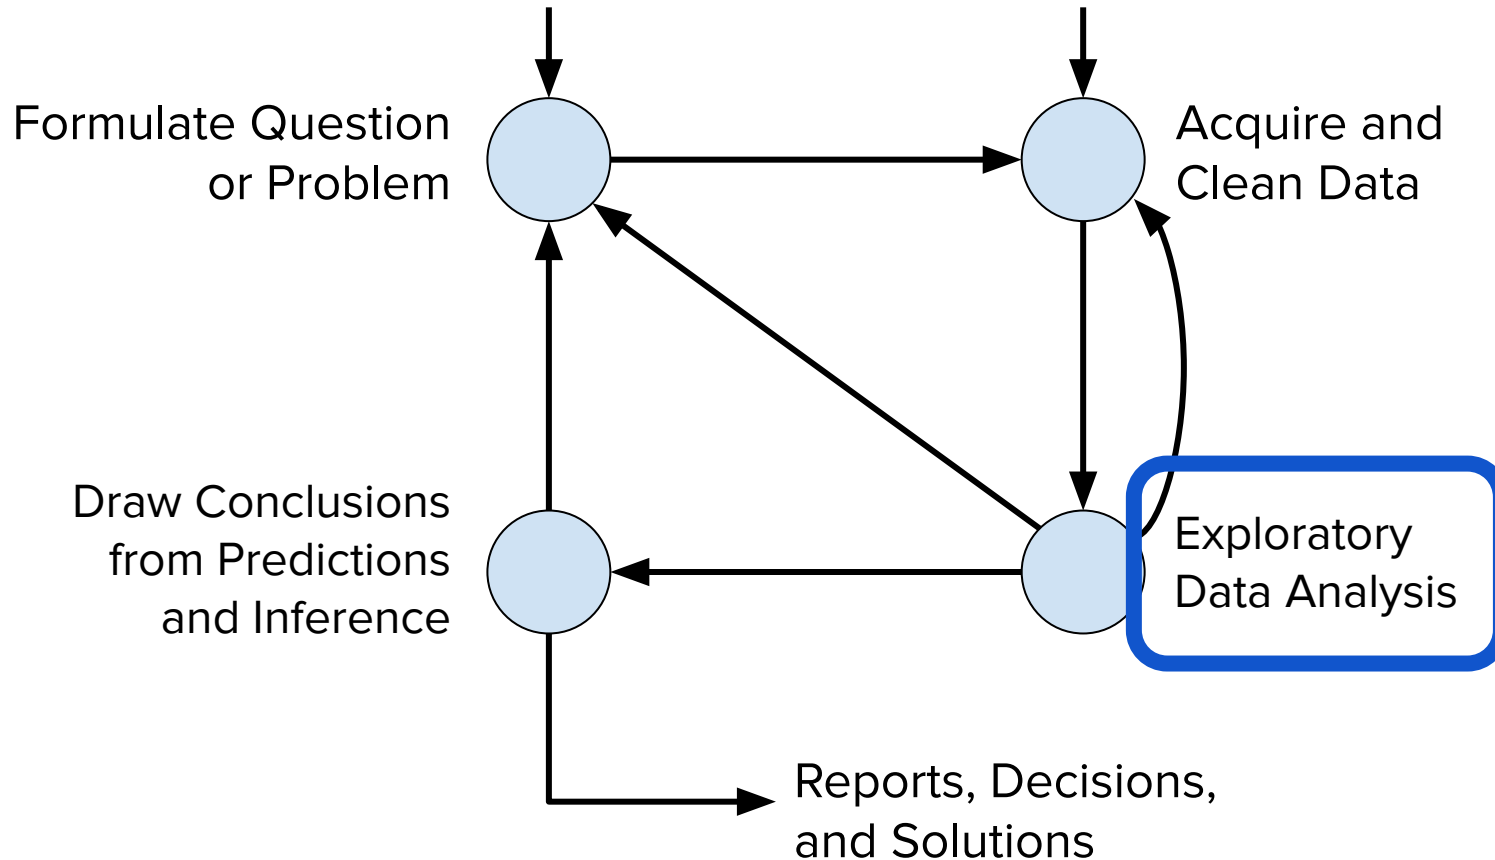
(Slides adapted from John DeNero, Josh Hug, Joey Gonzalez, and Fernando Perez)

Announcements

- HW2 out today
 - Will be “officially” due Friday but we will take submissions without penalty until Tuesday (July 9)
- HW3 out Friday
 - Due the following Friday (July 12)
- **For deadline extension requests, contact your TA and cc Sam in the email.**
- Congrats on finishing Project 1!
 - Will use Project 1 as case study today

Exploratory Data Analysis (EDA)

Remember the Data Science Lifecycle?



How Do We Understand the Data?

- Iterative process of visualizing and transforming data to:
 - Understand patterns
 - Identify issues
 - Inform subsequent analysis
 - Discover potential hypotheses (careful with this one!)
- Have an open-ended attitude
 - Look for what you believe is there and what you believe is not there (Tukey, 1965)

Goals of EDA

We want to determine the following:

Data Types What kinds of data do we have?

Granularity How fine/coarse is each datum?

Scope How (in)complete are the data?

Temporality How are the data situated in time?

Faithfulness How accurately do the data describe the world?

Data Types

Data Types: What kinds of data do we have?

- In Data 100:
 - Nominal Data: categories without natural ordering
 - Ordinal Data: categories with natural ordering
 - Numerical Data: amounts or quantities
- Note that these are **statistical data types**
 - Since these are more common, just say **data types**
 - These are distinct from **computational data types...**

Statistical vs. Computational Data Types

- Computational data types: int, float, string, boolean, etc.
- Each column in pandas has a computational data type
 - But does not imply a statistical data type
- Example: Sex sometimes coded as 0 for male, 1 for female, 2 for missing
 - Will be number in pandas, but should not be treated as numeric data!

Condensed Dataset from Project 1:

Computational

```
>>> df.dtypes
```

```
id          int64
lat         float64
phone       object
dtype: object
```

(“object” in
pandas usually
means string)

id

lat

phone

66 37.76 +14152427970

1085 37.79 +14152410256

1103 37.76 +14152410256

1116 37.80 +14159214049

1122 37.80 +14156736781

1127 37.79 +14159551940

1265 37.76 +14158540888

Statistical

What do you
think?

Condensed Dataset from Project 1:

Computational

```
>>> df.dtypes
```

```
id          int64
lat         float64
phone       object
dtype: object
```

(“object” in
pandas usually
means string)

id

lat

phone

66 37.76 +14152427970

1085 37.79 +14152410256

1103 37.76 +14152410256

1116 37.80 +14159214049

1122 37.80 +14156736781

1127 37.79 +14159551940

1265 37.76 +14158540888

Statistical

id: nominal

lat: numeric

phone: nominal

Granularity

Granularity: What does each record represent?

Do all records capture granularity at the same level?

- Some data will include summaries as records

If the data are coarse, how were they aggregated?

- Sampling, averaging, etc.

What kinds of aggregation is possible?

- From individuals to categories
- From individual events to totals across time or space

How might we adjust the granularity?

Project 1 Question 2b

- What does each record represent (e.g., a business, a restaurant, a location, etc.)?
- What is the primary key?
- What would you find by grouping by the following columns: `business_id`, `name`, `address` each individually?

	business_id	name	address	city	...	latitude	longitude	phone_number
0	19	NRGIZE LIFESTYLE CAFE	1200 VAN NESS AVE, 3RD FLOOR	San Francisco	...	37.79	-122.42	+14157763262
1	24	OMNI S.F. HOTEL - 2ND FLOOR PANTRY	500 CALIFORNIA ST, 2ND FLOOR	San Francisco	...	37.79	-122.40	+14156779494
2	31	NORMAN'S ICE CREAM AND FREEZES	2801 LEAVENWORTH ST	San Francisco	...	37.81	-122.42	NaN
...
6403	94571	THE PHOENIX PASTIFICIO	200 CLEMENT ST	San Francisco	...	NaN	NaN	+14154726100
6404	94572	BROADWAY DIM SUM CAFE	684 BROADWAY ST	San Francisco	...	NaN	NaN	NaN
6405	94574	BINKA BITES	2241 GEARY BLVD	San Francisco	...	NaN	NaN	+14157712907

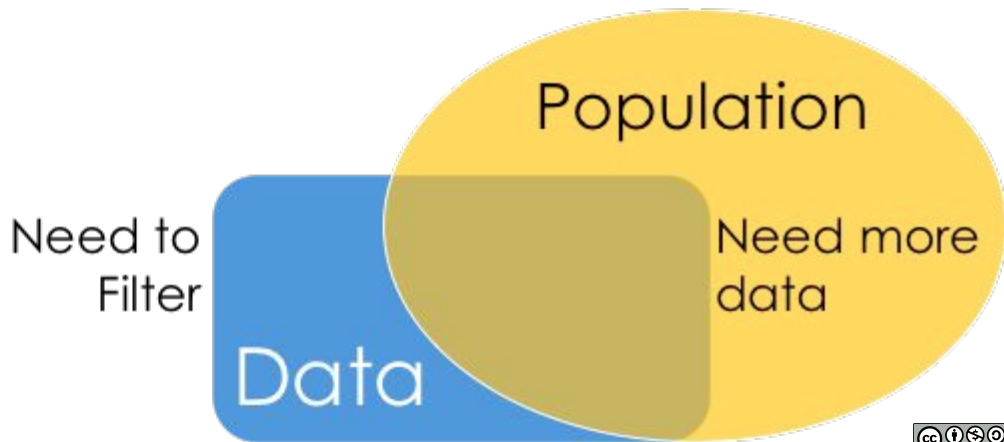
Scope

Scope: What do the data describe?

Do my data cover my area of interest?

Example: I am interested in studying crime in California, but I only have Berkeley + Seattle crime data.

In general, can only make statistical inferences about **probability samples** from **population of interest**!



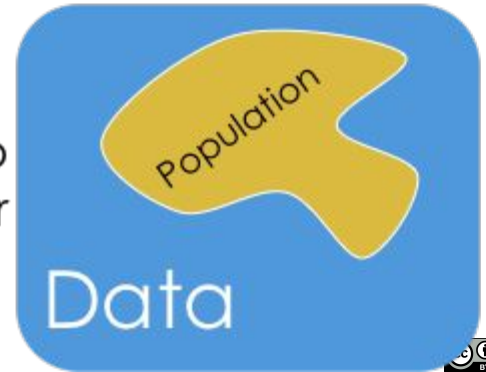
Scope: What do the data describe?

Do my data cover more than my area of interest?

Example: I am interested in studying Berkeley, but I have data for all of CA.

A sample over a large scope may have poor coverage for a narrow scope.

Need to
Filter



Scope: What do the data describe?

What about the businesses dataset in project 1?

- Is it a sample? Or census?
- What time frame?

Temporality

Temporality: When were the data collected?

What is the meaning of each date/time field?

- When the described event happened?
- When the data were collected or entered into the system?
- When the data were copied?

Temporality: When were the data collected?

Time and time formatting depend on location!

- Time zones
- Daylight savings (which transitions on different days in different regions)
- Regional formatting (06/07 vs 07/06)

Temporality: When were the data collected?

Are there strange zero/null values?

- In Unix / POSIX systems, time is measured from January 1st 1970 Coordinated Universal Time (UTC): 0 degrees latitude w/o daylight savings
- Be wary if you see a lot of 1969 or 1970 timestamps

Faithfulness

Faithfulness: Should I trust the data?

Do the data contain unrealistic or incorrect values?

- Dates in the future for events in the past
- Locations that don't exist
- Negative counts
- Misspellings of names
- Incorrectly typed numbers that are out of scale

Faithfulness: Should I trust the data?

Do the data violate obvious dependencies?

- E.g., age and birthday don't match

Were the data entered by hand?

- E.g., spelling errors, fields shifted

Did the data entry form provide default values?

Are there signs of deliberate data falsification?

- Repeated names, fake email addresses, repeated use of uncommon names.

EDA and Data Cleaning are an Iterative Cycle

Project 1 Zip Codes

Let's plot the zip codes.

Missing values! Let's just drop for now and replot.

```
bus[ 'postal_code' ].astype(int).hist()
```

```
~/anaconda3/lib/python3.7/site-packages/pandas/c
681         # work around NumPy brokenness,
682         if np.issubdtype(dtype.type, np.
--> 683             return lib.astype_intsafe(ar
684
685         # if we have a datetime/timedelt
```

```
pandas/_libs/lib.pyx in pandas._libs.lib.astype_
```

```
ValueError: cannot convert float NaN to integer
```

Project 1 Zip Codes

Weird zip codes! We'll only keep zip codes in SF.

```
bus['postal_code'].dropna().astype(int).hist()
```

ValueError

Traceback (most recent call last)

```
~/anaconda3/lib/python3.7/site-packages/pandas/core/dtypes/cast.py
    681         # work around NumPy brokenness, #1987
    682         if np.issubdtype(dtype.type, np.integer):
--> 683             return lib.astype_intsafe(arr.ravel(), dtype).r
    684
    685         # if we have a datetime/timedelta array of objects
```

```
pandas/_libs/lib.pyx in pandas._libs.lib.astype_intsafe()
```

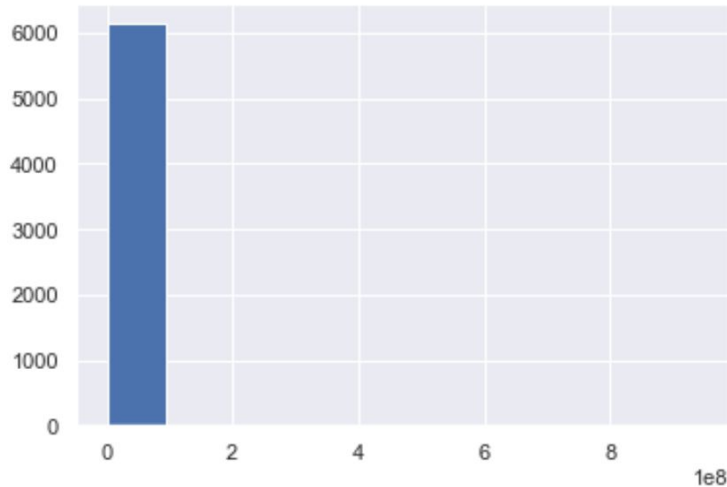
ValueError: invalid literal for int() with base 10: 'CA'

Project 1 Zip Codes

Was this what we expected for a histogram of zip codes?

```
(bus.loc[bus['postal_code_5'].isin(all_sf_zip_codes), 'postal_code']  
  .dropna()  
  .astype(int)  
  .hist()  
)
```

<matplotlib.axes._subplots.AxesSubplot at 0x1a2b60f278>



Project 1 Zip Codes

Looks like there are
malformed zip codes!

We'll clean those to be
five digits.

```
bus['postal_code'].value_counts()
```

94110	596
94103	552
94102	462

941033148	1
94102	1
941102019	1

Name: postal_code, Length: 33, dtype: int64

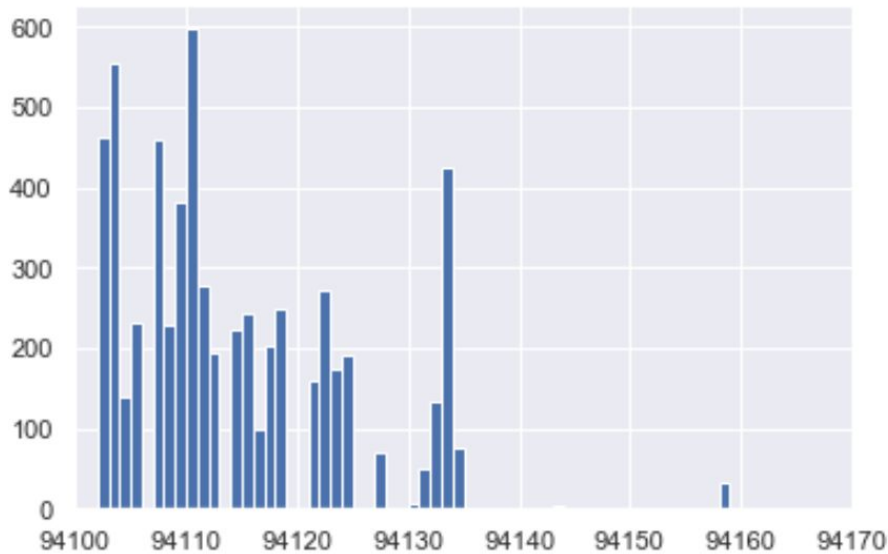
Project 1 Zip Codes

Now our zip codes look reasonable.

Now we turn back to missing zip codes, etc.

```
(bus
 .assign(postal_code_5=bus['postal_code'].str[:5])
 .loc[bus['postal_code_5'].isin(all_sf_zip_codes), 'postal_code_5']
 .dropna()
 .astype(int)
 .hist(bins=np.arange(94100, 94170))
)
plt.xlim(94100, 94170)
```

(94100, 94170)



Break!

Fill out Attendance:

<http://bit.ly/at-d100>

Text Fields

Text Fields and Data Cleaning / EDA

Extract quantitative values from text: dates, times, positions, etc.

Determine if missing values are denoted with a special string:
try `value_counts`

```
169.237.46.168 - - [26/Jan/2014:10:47:58  
-0800] "GET /stat141/Winter04/ HTTP/1.1" 200  
2585 "http://anson.ucdavis.edu/courses/"
```

(demo)

Regular Expressions

Regular Expressions

A *formal language* is a set of strings, typically described implicitly.

A *regular language* is a language that can be described by a *regular expression*.

- What's “regular” about it? Repeated and optional parts of a string, but no nested structures such as checking for balanced parentheses.
- In practice, many regular expression implementations can recognize more than just regular languages, including nested structures.

Regular Expressions

Example: `[0-9]{3}-[0-9]{2}-[0-9]{4}`

3 of any digit, then a dash, then 2 of any digit, then a dash, then 4 of any digit.

```
text = "My social security number is 123-45-6789.";
pattern = "[0-9]{3}-[0-9]{2}-[0-9]{4}"
re.findall(pattern, text)
```

Regular Expression Syntax

The four basic operations for regular expressions.

operation	order	example	matches	does not match
concatenation	3	AABAAB	AABAAB	any other string
or	4	AA BAAB	AA BAAB	any other string
closure (zero or more)	2	AB*A	AA ABBBBBBA	AB ABABA
parentheses	1	A(A B)AAB	AAAAB ABAAB	any other string
		(AB)*A	A ABABABABA	AA ABBA

Regular Expression Syntax

AB^* : A then zero or more copies of B: A, AB, ABB, AB BB

$(AB)^*$: Zero or more copies of AB: ABABABAB, ABAB, AB

operation	order	example	matches	does not match
concatenation	3	AABAAB	AABAAB	any other string
or	4	$AA BAAB$	AA BAAB	any other string
closure (zero or more)	2	AB^*A	AA ABBBBBBA	AB ABABA
parentheses	1	$A(A B)AAB$	AAAAB ABAAB	any other string
		$(AB)^*A$	A ABABABABA	AA ABBA

Expanded Regex Syntax

operation	example	matches	does not match
any character (except newline)	.U.U.U.	CUMULUS JUGULUM	SUCCUBUS TUMULTUOUS
character class	[A-Za-z][a-z]*	word Capitalized	camelCase 4illegal
at least one	jo+hn	john joooooooohn	jhn jjohn
zero or one	joh?n	jon john	any other string
repeated exactly {a} times	j[aeiou]{3}hn	jaoehn joohn	jhn jaeiouhn
repeated from a to b times: {a,b}	j[ou]{1,2}hn	john juohn	jhn joohn

REGULAR EXPRESSION v1 ▾

1 match, 35 steps (~0ms)

`"[0-9]{3}-[0-9]{2}-[0-9]{4}"`

TEST STRING

SWITCH TO UNIT TESTS ▸

My social security number is 123-45-6789.

EXPLANATION ▾

- ▼ `"[0-9]{3}-[0-9]{2}-[0-9]{4}"`
 - ▼ Match a single character present in the list below
 - `[0-9]{3}`
{3} Quantifier — Matches exactly 3 times
`0-9` a single character in the range between `0` (index 48) and `9` (index 57) (case

MATCH INFORMATION ▾

Match 1

Full match 29-40 123-45-6789

Puzzle: Use regex101.com to test!

<https://www.debuggex.com/cheatsheet/regex/python>

Regular Expression Basics	
.	Any character except newline
a	The character a
ab	The string ab
a b	a or b
a*	0 or more a's

Regular Expression Quantifiers	
*	0 or more
+	1 or more
?	0 or 1
{2}	Exactly 2
{2, 5}	Between 2 and 5

Give a regular expression that matches “moon”, “mooon”, “moooooon”, etc.

Your expression should match any **positive, even** number of o's.

Solution: <https://regex101.com/r/dj0nqD/1>

Puzzle: Use regex101.com to test!

<https://www.debuggex.com/cheatsheet/regex/python>

Regular Expression Basics

.	Any character except newline
a	The character a
ab	The string ab
a b	a or b
a*	0 or more a's

Regular Expression Quantifiers

*	0 or more
+	1 or more
?	0 or 1
{2}	Exactly 2
{2, 5}	Between 2 and 5

Give a regex that matches muun, muuuun, moon, moooon, etc. Your expression should match any positive, even number of either u's or o's, but not a u/o mix.

Solution: <https://regex101.com/r/4CMWjj/1>

Order of Operations in Regexes

$m((uu)^+|(oo)^+)n$

- Matches starting with m and ending with n, with either of the following in the middle:
 - $(uu)^+$
 - $(oo)^+$

Match examples:

muun
muuuun
moon
moooon

$m((uu)^+)((oo)^+)n$

- Matches either of the following
 - m followed by $(uu)^+$
 - $(oo)^+$ followed by n

Match examples:

muu
muuuu
oon
ooooon

Limitations of Regular Expressions

Writing regular expressions is like writing a program.

- Need to know the syntax well.
- Can be easier to write than to read.
- Can be difficult to debug.
- Regular expressions are sometimes jokingly referred to as a [write only language](#).
- For parsing a hierarchical structure, such as JSON, use a parser, not a regex.

"Some people, when confronted with a problem, think 'I know, I'll use regular expressions.' Now they have two problems." - Jamie Zawinski ([Source](#))

Email Address Regular Expression

The regular expression for email addresses (for the Perl programming language):

[illegible]

From: <http://www.ex-parrot.com/~pdw/Mail-RFC822-Address.html>



Even More Regular Expression Syntax

operation	example	matches	does not match
built-in character classes	<code>\w+</code> <code>\d+</code>	fawef 231231	this person 423 people
character class negation	<code>[^a-z]+</code>	PEPPERS3982 17211!↑å	porch CLAmS
escape character	<code>cow\.com</code>	cow.com	cowscom

Suppose you want to match one of the reserved characters, such as `.` or `[` or `]`

- In these cases, you must *escape* the character using a backslash.

Regular Expressions Puzzle: tinyurl.com/reg913a

operation	example	matches	does not match
built-in character classes	<code>\w+</code> <code>\d+</code>	fawef 231231	this person 423 people
character class negation	<code>[^a-z]+</code>	PEPPERS3982 17211!↑å	porch CLAmS
escape character	<code>cow\.com</code>	cow.com	cowscom

Create a regular expression that matches the red portion below.

```
169.237.46.168 - - [26/Jan/2014:10:47:58  
-0800] "GET /stat141/Winter04/ HTTP/1.1" 200  
2585 "http://anson.ucdavis.edu/courses/"
```


Even More Regular Expression Features

operation	example	matches	does not match
beginning of line	<code>^ark</code>	ark two ark o ark	dark
end of line	<code>ark\$</code>	dark ark o ark	ark two
non-greedy qualifier	<code>5.*?5</code>	5005 55	5005005

A few additional common regex features are listed above.

5*5 would
match this!

- Won't discuss these in class, but might come up in discussion or hw.
- There are even more out there!

The official guide is good! <https://docs.python.org/3/howto/regex.html>

Matching Regular Expression Patterns in Python

re.findall in Python

In Python, `re.findall(pattern, text)` will return a list of all matches.

```
text = "My SSN is 456-76-4295, or maybe 456-67-4295."  
pattern = "[0-9]{3}-[0-9]{2}-[0-9]{4}"  
m = re.findall(pattern, text)  
print(m)
```

```
['456-76-4295', '456-67-4295']
```

Regular Expression Groups

Earlier we used parentheses to specify the order of operations.

Parentheses have another meaning:

- Every set of parentheses in the pattern corresponds to a group
- Regular expression matchers (e.g. `re.findall`, regex101.com) will return matches organized by groups.

```
s = """Observations: 03:04:53 - Horse awakens.  
03:05:14 - Horse goes back to sleep."""  
pattern = "(\d\d):(\d\d):(\d\d) - (.*)"   
matches = re.findall(pattern, s)
```

```
[('03', '04', '53', 'Horse awakens.'),  
 ('03', '05', '14', 'Horse goes back to sleep.)]
```

re.sub in Python

In Python, `re.sub(pattern, repl, text)` will return `text` with all instances of `pattern` replaced by `repl`.

```
text = '<div><td valign="top">Moo</td></div>'
pattern = "<[^>]+>"
cleaned = re.sub(pattern, '', text)
print(cleaned)
```

```
'Moo'
```

Optional (but Handy) Regex Concepts

These regex features aren't going to be on an exam, but they are useful:

- **Lookaround:** match “good” if it’s not preceded by “not”:

`(?<!not)good`

- **Backreferences:** match HTML tags of the same name:

`<(\w+)>.*</\1>`

Optional (but Handy) Regex Concepts

These regex features aren't going to be on an exam, but they are useful:

- **Named groups:** match a vowel as a named group:
`(?P<vowel>[aeiou])`
- **Free Space:** Allow free space and comments in a pattern:

Match a 20th or 21st century date in yyyy-mm-dd format

`(19|20)\d\d` *# year (group 1)*

`[- /.]` *# separator*

`(0[1-9]|1[012])` *# month (group 2)*

`[- /.]` *# separator*

`(0[1-9]|1[12][0-9]|3[01])` *# day (group 3)*

Use Case: Text Features

Top 20 Violations

unclean or degraded floors walls or ceilings

moderate risk food holding temperature

inadequate and inaccessible **handwashing** facilities

unapproved or unmaintained equipment or utensils

inadequately cleaned or **sanitized** food contact surfaces

wiping cloths not **clean** or properly stored or inadequate sanitizer

improper food storage

foods not protected from contamination

moderate risk **vermin infestation**

high risk food holding temperature

unclean nonfood contact surfaces

food safety certificate or food handler card not available

unclean or **unsanitary** food contact surfaces

inadequate food safety knowledge or lack of certified food safety manager

improper storage of equipment utensils or linens

low risk **vermin infestation**

permit license or inspection report not posted

improper cooling methods

unclean hands or improper use of **gloves**

improper or defective plumbing

What Features to Examine?

Cleanliness 'clean|sanit'

High risk 'high risk'

Vermin 'vermin'

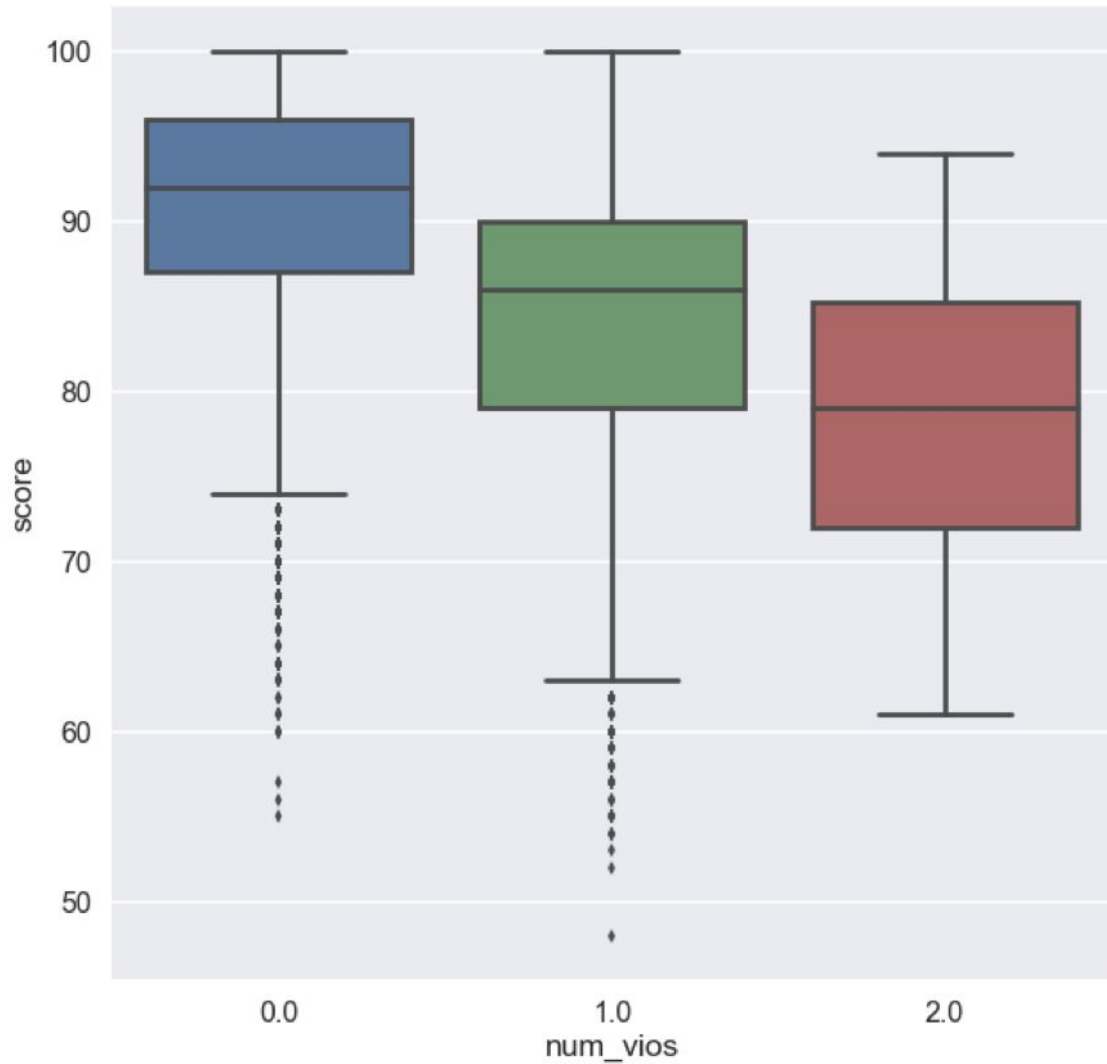
Surfaces 'wall|ceiling|floor|surface'

Humans 'hand|glove|hair|nail'

Permits and Certification 'permit|certif'

Vermin Violations

Median safety score is 86 for restaurants with one vermin infestation violation!



Summary

- EDA is a process used to understand a dataset
 - Investigate Data Types, Granularity, Scope, Temporality, and Faithfulness
- To extract features from text fields:
 - Use built-in Python string methods
 - Or regex if your pattern is complicated