

# Linear Regression

## (Reading: [Ch 13](#))

---

**UC Berkeley Data 100 Summer 2019**  
**Sam Lau**

### Learning goals:

- Reframe loss minimization framework for modeling.
- Introduce multivariable linear models within the loss minimization framework.

(Slides adapted from Sandrine Dudoit and Joey Gonzalez)

# Announcements

- HW4 due **Tuesday**
- HW5 out Tuesday, due **Friday**
- Leo was out sick! Hopefully back today.
- Today's lecture might get split into two
  - Ask lots of questions if we're going too fast

# Last Time

- Draw conclusions about population using a sample through statistical estimation.
- Make estimations by picking the estimator that minimizes empirical risk / loss.
- Today:
  - Connect estimation with prediction and modeling.
  - First foray into machine learning with linear models.

# Modeling


# Making Predictions

- Loss minimization framework useful for predictions too!
- Suppose we have a dataset of cars and we'd like to predict fuel efficiency (miles per gallon, or mpg):


	mpg	cylinders	displacement	horsepower	...	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	...	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	...	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	...	11.0	70	usa	plymouth satellite

# Models

- To make a prediction, we choose a **model**.
  - Takes input data and outputs a prediction.
- Constant model:



- Simple linear model:



# The Constant Model

- Start simple: if constant model, how do we pick  $\theta$ ?

$$f_{\theta}(x) = \theta$$

- Intuition: pick  $\theta$  to be close to most of the values in data

	mpg	cylinders	displacement	horsepower	...	acceleration	model_year	origin	name
0	18.0	8	307.0	130.0	...	12.0	70	usa	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	...	11.5	70	usa	buick skylark 320
2	18.0	8	318.0	150.0	...	11.0	70	usa	plymouth satellite

# Model Loss

- Use  $x_i$  to denote what we use to make predictions
- Use  $y_i$  to denote what we're trying to predict
- But both  $x$  and  $y$  come from a single sample
- Idea: Pick the  $\theta$  that minimizes the average loss between  $y$  in our sample and model predictions.

$$L(\theta, y_1, \dots, y_n) = \frac{1}{n} \sum (y_i - f_{\theta}(x))^2$$



# Constant Model Loss

$$L(\theta, y_1, \dots, y_n) = \frac{1}{n} \sum (y_i - f_\theta(x))^2$$

Since  $f_\theta(x) = \theta$  for constant model:

$$L(\theta, y_1, \dots, y_n) = \frac{1}{n} \sum (y_i - \theta)^2$$

- Remember this expression from last lecture?
- $\theta$  = sample mean is the best model parameter.
- So, for car MPGs, we set  $\theta = \text{mean(mpg)}$

# Modeling is Estimation in New Clothes

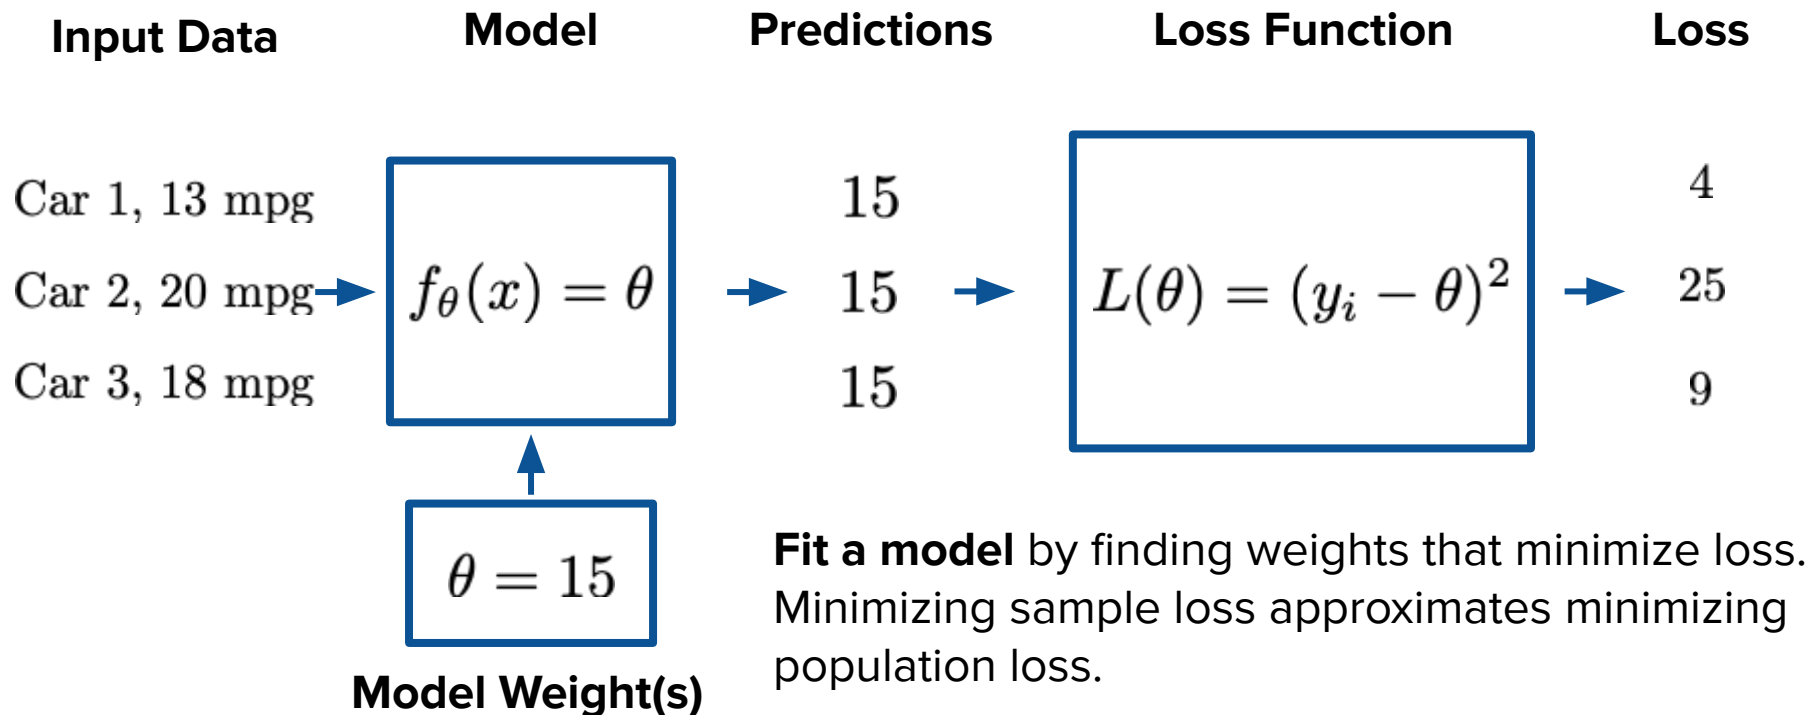
- Estimation: making best guess at population parameter
- Modeling: making predictions for population values
- Two sides of the same coin! Why?
- Modeling assumes pop values generated by parameters:

$$f_{\theta}^*(x) = \theta^* + \epsilon$$

- RTA: assume that data from population generated by taking a constant  $\theta^*$  and adding noise  $\epsilon$ .
- Estimation = Finding  $\hat{\theta}$ , our best estimate for  $\theta^*$
- Modeling = Using  $\hat{\theta}$  to make predictions

# The Modeling Pipeline

We choose what goes in the blue boxes!



# The Modeling Recipe

- Pick a model, pick a loss function, fit the model to sample.
- Preview of model and loss function combos:


Model	Loss Function	Technique Name
$\theta \cdot x$	$\frac{1}{n} \sum (y_i - f_{\theta}(x))^2$	Least squares linear regression
$\theta \cdot x$	$\frac{1}{n} \sum (y_i - f_{\theta}(x))^2 + \lambda  \theta ^2$	Lasso regression
$\theta \cdot x$	$\frac{1}{n} \sum (y_i - f_{\theta}(x))^2 + \lambda \ \theta\ _{\ell_1}$	Ridge regression
$\theta \cdot x$	$\frac{1}{n} \sum  y_i - f_{\theta}(x) $	Least absolute deviations
$\sigma(\theta \cdot x)$	$\frac{1}{n} \sum [-y_i \ln f_{\theta}(x) - (1 - y_i) \ln(1 - f_{\theta}(x))]$	Logistic regression

# Linear Models

# Using Our Data

- If we're trying to predict MPG, we can do better than a constant model by incorporating more information.
  - E.g. higher horsepowers have lower MPGs:

	mpg	cylinders	displacement	horsepower
0	18.0	8	307.0	130.0
1	15.0	8	350.0	165.0
2	18.0	8	318.0	150.0



# Simple Linear Model

- We want our predictions to depend on the input data  $x$ .
- Simple linear model:

$$f_{\theta}^*(x) = \theta_1^* x + \theta_0^* + \epsilon \quad \epsilon \sim N(0, \sigma^2)$$

$$f_{\theta}(x) = \theta_1 x + \theta_0$$

- As usual, we can minimize the loss. This time, we have two parameters.

$$\begin{aligned} L(\theta_1, \theta_0, y_1, \dots, y_n) &= \frac{1}{n} \sum (y_i - f_{\theta}(x_i))^2 \\ &= \frac{1}{n} \sum (y_i - \theta_1 x_i - \theta_0)^2 \end{aligned}$$

# Simple Linear Model

$$L(\theta_1, \theta_0, y_1, \dots, y_n) = \frac{1}{n} \sum (y_i - \theta_1 x_i - \theta_0)^2$$

$$\frac{\partial}{\partial \theta_1} L = \frac{1}{n} \sum 2(y_i - \theta_1 x_i - \theta_0)(-x_i)$$

$$\frac{\partial}{\partial \theta_0} L = \frac{1}{n} \sum 2(y_i - \theta_1 x_i - \theta_0)(-1)$$

- This ends up being a lot of algebra, so we'll skip to the answer.



# Skipping Ahead

Let  $r$  be the average of the products of  $x$  and  $y$  when both variables are measured in standard units.

$$\hat{\theta}_1 = r \cdot \frac{\text{SD of } y}{\text{SD of } x}$$

$$\hat{\theta}_0 = \text{mean of } y - \hat{\theta}_1 \cdot \text{mean of } x$$

- [Data 8 textbook](#) has example slope/intercept calculations.
- Takeaway: Can derive these formulas by minimizing loss.
- You should know how to take the derivative but won't need to solve it.

# Multivariable Linear Model

- Simple linear model uses one variable to predict:

$$f_{\theta}(x) = \theta_1 x + \theta_0$$

- Time to graduate from Data 8!
- Multivariable linear model uses  $\geq 1$  variable:



$$f_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_p x_p$$

- $\mathbf{x}$  is a vector containing one row of input data.
- IOW: Predict by combining multiple features together.

# Intuition

$$f_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_p x_p$$

- Using horsepower and model year to predict mpg
  - Expect  $\theta_1$  to be negative and  $\theta_2$  to be positive. Why?



# Using Matrix Multiplication

$$f_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_p x_p$$

- Many terms to write! We'll use a trick: add a column of 1s to the table:

	<b>mpg</b>	<b>bias</b>	<b>horsepower</b>	<b>weight</b>	<b>model_year</b>
<b>0</b>	18.0	1	130.0	3504	70
<b>1</b>	15.0	1	165.0	3693	70
<b>2</b>	18.0	1	150.0	3436	70

$$\mathbf{x} = [1, x_1, x_2, x_3]$$

$$\boldsymbol{\theta} = [\theta_0, \theta_1, \theta_2, \theta_3]$$

Bolded letters means vector or matrix.

- This means our model is:  $f_{\theta}(\mathbf{x}) = \boldsymbol{\theta} \cdot \mathbf{x}$

# More Notation!

$y$ : column vector of sample points to predict

$X$ : matrix of input data ( $n \times p$ )

	mpg	bias	horsepower	weight	model_year
0	18.0	1	130.0	3504	70
1	15.0	1	165.0	3693	70
2	18.0	1	150.0	3436	70
...	...	...	...	...	...
395	32.0	1	84.0	2295	82
396	28.0	1	79.0	2625	82
397	31.0	1	82.0	2720	82

$X_i$ : row  $i$  of input data

$x$ : row vector of input data

$\theta$ : vector of model weights

$\hat{\theta}$ : loss-minimizing model weights

**Your turn:** Write the matrix expression that computes a vector with a fitted linear model's predictions for **all sample points**.

## Your Turn

Write the matrix expression that computes a vector with a fitted linear model's predictions for all sample points.

Prediction for one point:  $\hat{\theta} \cdot x$

$$\text{Prediction for all points: } \hat{y} = \begin{bmatrix} \hat{\theta} \cdot X_1 \\ \hat{\theta} \cdot X_2 \\ \vdots \\ \hat{\theta} \cdot X_n \end{bmatrix} = X\hat{\theta}$$

$(n \times p)(p \times 1) = (n \times 1)$

## Your Turn

Write the matrix expression that computes the average MSE loss for all data points (this is a scalar!).

## Your Turn

Write the matrix expression that computes the average MSE loss for all data points (this is a scalar!).

$$L(\boldsymbol{\theta}, \mathbf{X}, \mathbf{y}) = \frac{1}{n} \sum (y_i - \mathbf{X}_i \cdot \boldsymbol{\theta})^2$$

$$= \frac{1}{n} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|^2$$

(Bonus points if you got this)

$$\|\mathbf{v}\|^2 = \mathbf{v} \cdot \mathbf{v} = \mathbf{v}^\top \mathbf{v}$$

Using matrix notation takes a lot of practice to get used to, but the results are worth it. Always check your dimensions!



# Fitting a Linear Model

- How do we pick  $\theta$  to minimize loss?

$$L(\theta, \mathbf{X}, \mathbf{y}) = \frac{1}{n} \sum (y_i - \mathbf{X}_i \cdot \theta)^2$$

- Want to take partial derivatives for  $\theta_0, \theta_1, \dots$
- Instead, we'll take the **gradient** and set it equal to zero.
- This solves for all model weights at once!

$$\nabla_{\theta} L = \begin{bmatrix} \frac{\partial}{\partial \theta_0} (L) \\ \frac{\partial}{\partial \theta_1} (L) \\ \vdots \\ \frac{\partial}{\partial \theta_p} (L) \end{bmatrix}$$

# The Normal Equation

- Saving the setup for the Gradient Descent lecture
  - Again, you need to know how to take the gradient but not how to solve for  $\theta$ .
- Skipping ahead to the answer:

$$\mathbf{X}^\top \mathbf{X} \hat{\theta} = \mathbf{X}^\top \mathbf{y}$$

$$\hat{\theta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

What are the matrix shapes in these expressions?

- Expression above called normal equation
- Gives a closed-form recipe for fitting linear model

# The Abnormal Equation

- In practice, it takes too long to compute this:

$$\hat{\theta} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

- Inverting an (n x n) matrix takes at least  $O(n^2)$  time.
  - State of the art:  $O(n^{2.3})$
- Takeaway: analytic solutions are elegant but are sometimes hard to find and slow.
  - Next lecture: gradient descent

# Demo: Predicting MPGs

**Break!**

**Fill out Attendance:**


**<http://bit.ly/at-d100>**

# **Feature Engineering**

## **(moved to Wed lecture)**

# Linear Models Level Up

- Horsepower and mpg have a nonlinear relationship.
- Can still use linear regression to capture this!
- **Feature engineering**: creating new features from data to give model more complexity.



# Adding Features

- For now, predict MPG from horsepower alone.
- Insight: Add a new column to  $X$  with horsepower<sup>2</sup>.

	bias	hp	hp^2
0	1	130.0	16900.0
1	1	165.0	27225.0
2	1	150.0	22500.0
...	...	...	...
395	1	84.0	7056.0
396	1	79.0	6241.0
397	1	82.0	6724.0

- Now we fit a quadratic function!

$$\begin{aligned}f_{\theta}(\mathbf{x}) &= \theta_0 + \theta_1 x_1 + \theta_2 x_2 \\ &= \theta_0 + \theta_1 \text{hp} + \theta_2 \text{hp}^2\end{aligned}$$

- This is still linear in model weights  $\theta$ , so we call it a linear model.

(Demo)



# Polynomial Regression

- For polynomial features of degree  $n$ , usually add every possible combination of columns.
  - 4 original columns, degree 2:

$$x_1, x_2, x_3, x_4,$$

$$x_1^2, x_2^2, x_3^2, x_4^2,$$


$$x_1x_2, x_1x_3, x_1x_4, x_2x_3, x_2x_4, x_3x_4$$

- Can end up being a lot of columns
- To cope, use kernel trick (covered in advanced courses)

# Categorical Features

- Origin column is correlated with MPG. Can we use it?
- Idea: Encode categories as numbers in a smart way.
- Discuss: Why can't we just encode "usa" as 0, "japan" as 1, "europe" as 2?


```
sns.boxplot(x='origin', y='mpg', data=mpg);
```



# One-Hot Encoding

- One-hot encoding makes one new column for each unique category:

origin	
usa	
usa	
europa	
...	
usa	
japan	
japan	




origin=usa	origin=europa	origin=japan
1	0	0
1	0	0
0	1	0
...	...	...
1	0	0
0	0	1
0	0	1

# One-Hot Encoding

- What do you expect the largest weight to be?

origin=usa	origin=europe	origin=japan
1	0	0
1	0	0
0	1	0
...	...	...
1	0	0
0	0	1
0	0	1

```
sns.boxplot(x='origin', y='mpg', data=mpg);
```



- Can interpret weight as “contribution” of that category

# One Hot Problem

- Problem: Adding a new column for each category makes columns of  $X$  **linearly dependent**! Why?
- One-hot columns always sum to 1:

bias	origin=usa	origin=europe	origin=japan
1	1	0	0
1	1	0	0
1	0	1	0
...	=	+	+
1	1	0	0
1	0	0	1
1	0	0	1

- This makes normal equations unsolvable.

$$\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Not invertible ^

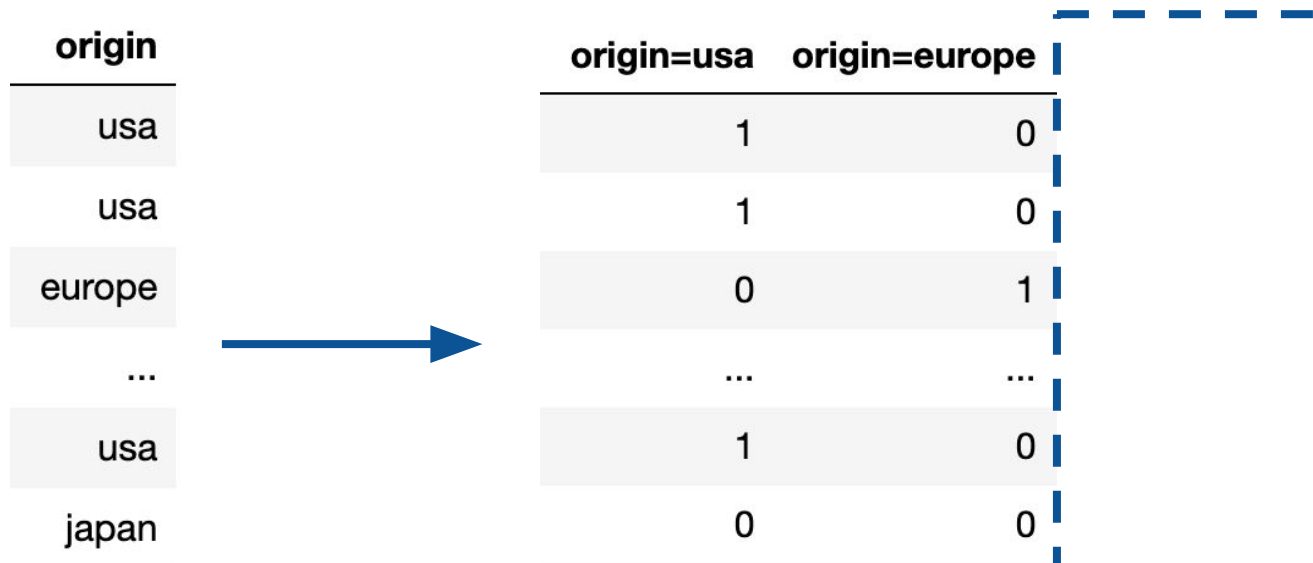
# Weight Interpretation

- Invertibility isn't a problem for gradient descent, but this still affects how we interpret the model weights.
- Linearly dependent columns can “swap” weights:
  - Left: All categories matter. Right: No categories matter!

0	3	3	3		3	0	0	0
bias	origin=usa	origin=europe	origin=japan		bias	origin=usa	origin=europe	origin=japan
1	1	0	0	=	1	1	0	0
1	1	0	0		1	1	0	0
1	0	1	0		1	0	1	0
...	...	...	...		...	...	...	...
1	1	0	0		1	1	0	0
1	0	0	1		1	0	0	1

# Drop it Like it's Hot

- Simple fix: Drop the last one-hot column.
- In this case, the weight for USA can be interpreted as “change in MPG between USA and Japan”.



origin		origin=usa	origin=europe
usa		1	0
usa		1	0
europe		0	1
...		...	...
usa		1	0
japan		0	0

# Features feat. More Features

- Feature engineering is often domain-specific:
  - Standardizing: “How many SDs away from average?”
  - Log transform: Used to fit exponential models.
  - Absolute difference: “How different is the current temperature from 70°?”
  - Binning data, then one-hot encoding: “Are we driving during morning rush hour? Evening rush hour?”
  - Date-related features: year, month, weekday
  - Image-related features: blurring, edge detection, etc.



# Summary

- Modeling and estimation are closely related.
  - We can view modeling as estimation of model parameters.
- Linear models can incorporate an arbitrary number of features to make a prediction.
- Feature engineering extends linear models to generate more complex models.